

Computer Networks

Chapter 3.5

Jong-won Lee

Handong University

Chapter 3

- ❑ 3.1 Transport-layer services
- ❑ 3.2 Multiplexing and demultiplexing
- ❑ 3.3 Connectionless transport: UDP
- ❑ 3.4 Principles of reliable data transfer (ARQ)
- ❑ 3.5 Connection-oriented transport: TCP
 - segment structure
 - RTT estimation and timeout
 - reliable data transfer
 - flow control
 - connection management
- ❑ 3.6 Principles of congestion control
- ❑ 3.7 TCP congestion control

TCP: Overview RFCs: 793, 1122, 1323, 2018, 2581

❑ connection-oriented:

- (exchange of control msgs) init's sender, receiver state before data exchange

❑ point-to-point:

❑ full duplex data:

- bi-directional data flow in same connection
- MSS:
 - The max. amount of application data in a segment

TCP: Overview RFCs: 793, 1122, 1323, 2018, 2581




- Sliding window protocol



□ Flow control

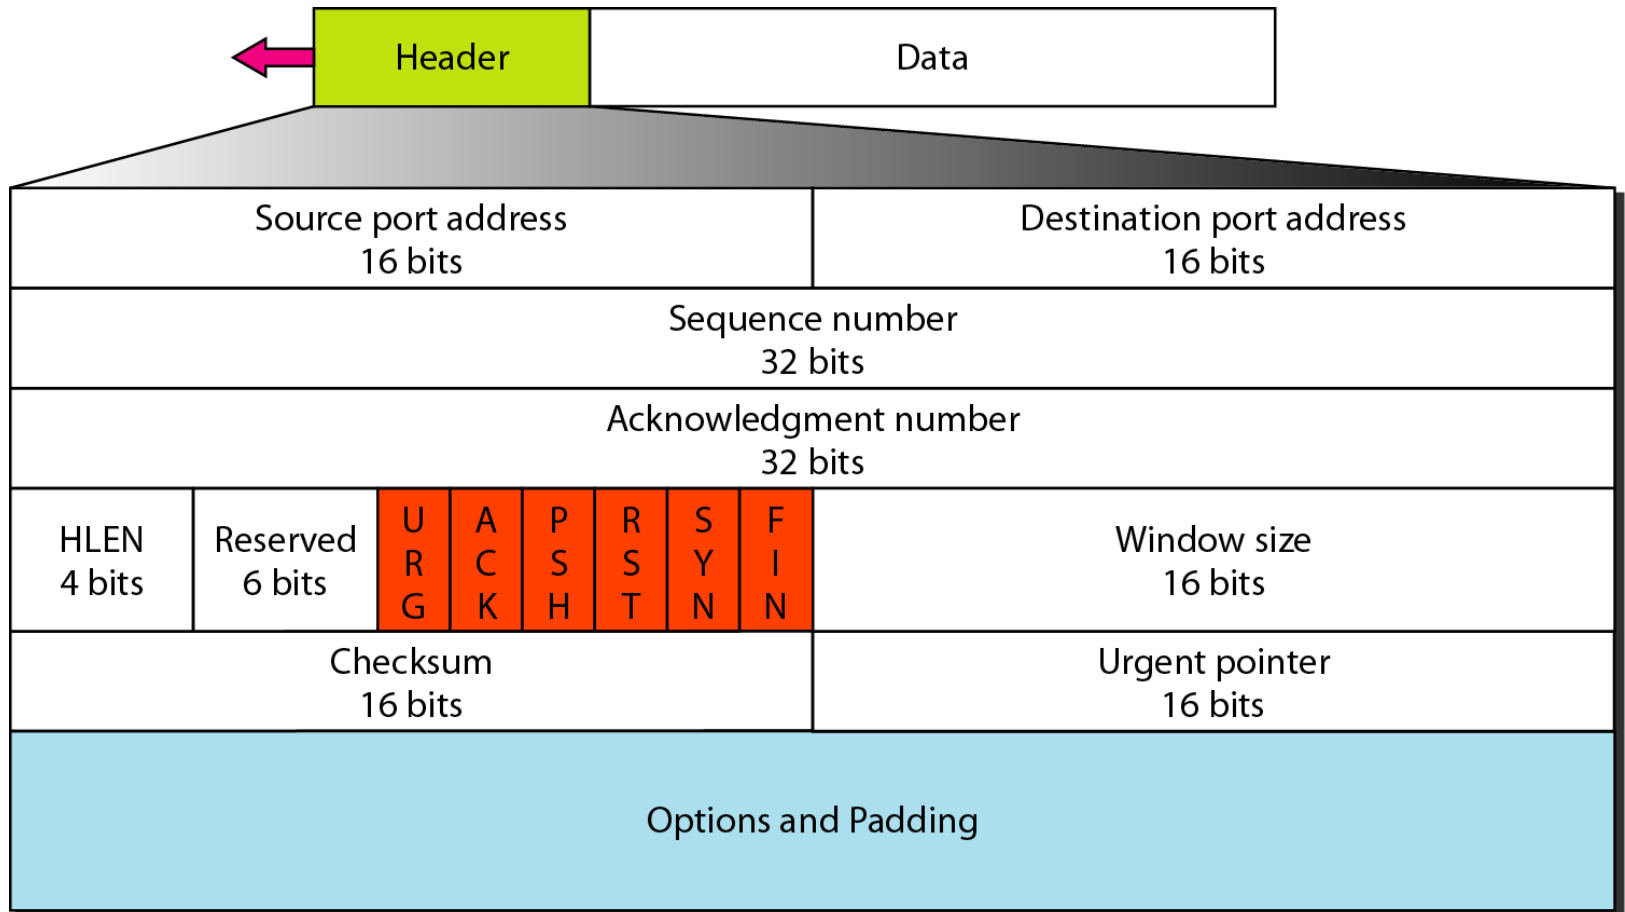
- sender will not overwhelm receiver buffer

□ Congestion control



TCP Segment Structure

- Header = 20 bytes + options



TCP Header

- ❑ *Source/destination port (16 bits)*
 - Identifies the application process
- ❑ *Sequence number (32 bits)*
 - Sequence number of the first byte in the segment.
 - Seq. number is advanced
 - If SYN is present, this is the initial sequence number (ISN) and the first data byte will be ISN+1.
 - during the connection establishment, each party uses a random number generator to create an *initial sequence number (ISN)*
- ❑ *Acknowledgement number (32 bits)*
 - Next expected sequence number, valid only when the ACK bit (reside in flag) is set
 - Cumulative ACK

TCP seq. #'s and ACKs

Seq. #'s:


- byte stream "number" of first byte in segment's data

ACKs:

- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- A: TCP spec doesn't say, - up to implementor



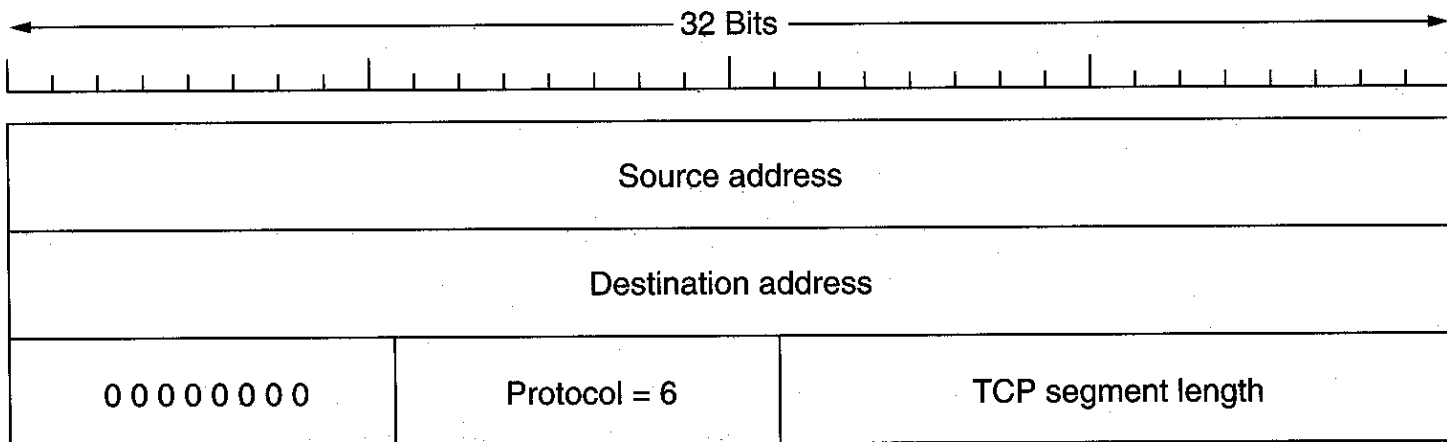
TCP Header

□ 6 control flag bits

- **URG** : set to 1 if *urgent pointer* is in use.
- : set to 1 if *Acknowledgement number* is valid
- **PSH** : set to 1 if the receiver should pass this data to the application as soon as possible.
- **RST** : set to 1 to reset a erroneous connection
- : synchronize sequence numbers to initiate a connection
 - If SYN=1 & ACK=0,
then this segment is a **CONNECTION-REQUEST**
 - If SYN=1 & ACK=1,
then this segment is a **CONNECTION-ACCEPTED**
- : set to 1 to indicate no more data to send

TCP Header

- ❑ *TCP header length (4 bits)*
 - the length of the header in 32-bit words
 - 4 bits => max header size is 60 bytes
 - needed because the options field is of variable length
- ❑ *Window size (16 bits)*
 - will accept [Ack] to [Ack]+[window-1]
- ❑ *Checksum*
 - Internet checksum of *the header , the data, and the pseudo-header*
 - The pseudo-header is not transmitted.



TCP Header

- ❑ *Urgent pointer (16 bits)*
 - valid only if the *URG* flag is set
 - Lets receiver know how much data it should deliver right away.
 - Points to **the last byte** of the urgent data. (urgent data is at the beginning of the segment)

- ❑ *Options (variable)*
 - MSS (Maximum Segment Size), window scale factor, timestamp etc.

RTT Estimation and Timeout

- Q: How to set TCP timeout value?
 - longer than RTT
 - but large variance of RTT
 - too short: premature timeout
 - unnecessary retransmissions
 - too long: slow reaction to segment loss
- A highly dynamic algorithm is needed
 - adjusts the timeout interval based on continuous measurements of network performance

RTT Estimation and Timeout


❑ Original RTT estimation method (RFC 793)

(Exponential weighted moving average)

- influence of past sample decreases exponentially fast
- EstimatedRTT =
 - (typically $\alpha = 7/8$)
 - **SampleRTT**: measured time from segment transmission until ACK receipt
- RTO =
- Too large RTO value

Example: RTT estimation:

RTT: gaia.cs.umass.edu to fantasia.eurecom.fr



RTT Estimation and Timeout

❑ Improved algorithm (Jacobson)

- β is roughly proportional to the standard deviation
- use a mean deviation D as an estimator of standard deviation

$$(\sum |M_i - a|)^2 \geq \sum |M_i - a|^2$$

- $\text{Err} = \text{SampleRTT} - \text{EstimatedRTT}$

$$\text{EstimatedRTT} = \text{EstimatedRTT} + g * \text{Err}$$

$$(\text{typically } g = 1/8 (=1-\alpha))$$

$$\text{DevRTT} = \text{DevRTT} + h * (|\text{Err}| - \text{DevRTT})$$

$$(\text{typically } h = 1/4)$$

$$\text{RTO} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$

RTT Estimation and Timeout

- Initial value for RTO:

- Sender should set the initial value of RTO to

$$RTO_0 = 3 \text{ seconds}$$

- RTO calculation after first RTT measurements arrived

$$\text{EstimatedRTT}_1 = \text{RTT}$$

$$\text{DevRTT}_1 = \text{RTT} / 2$$

$$RTO_1 = \text{EstimatedRTT}_1 + 4 * \text{DevRTT}_1$$

- When a timeout occurs, the RTO value is doubled

$$RTO_{n+1} = \min (q * RTO_n, RTO_{\max}) \text{ seconds}$$

(typically $q = 2$, $RTO_{\max} \geq 60 \text{ sec.}$)

TCP reliable data transfer

- ❑ Based on sliding window protocol
 - Similar to SR ARQ.
 - Dynamic window size
 - Combined with flow control & congestion control
 - Use cumulative ACK.
 - The value of the acknowledgment field is the number of the next byte that the receiver expects to receive.
 - Retransmissions are triggered by:
 - timeout events
 - Three duplicate ACKs (fast retransmission)

Fast Retransmit

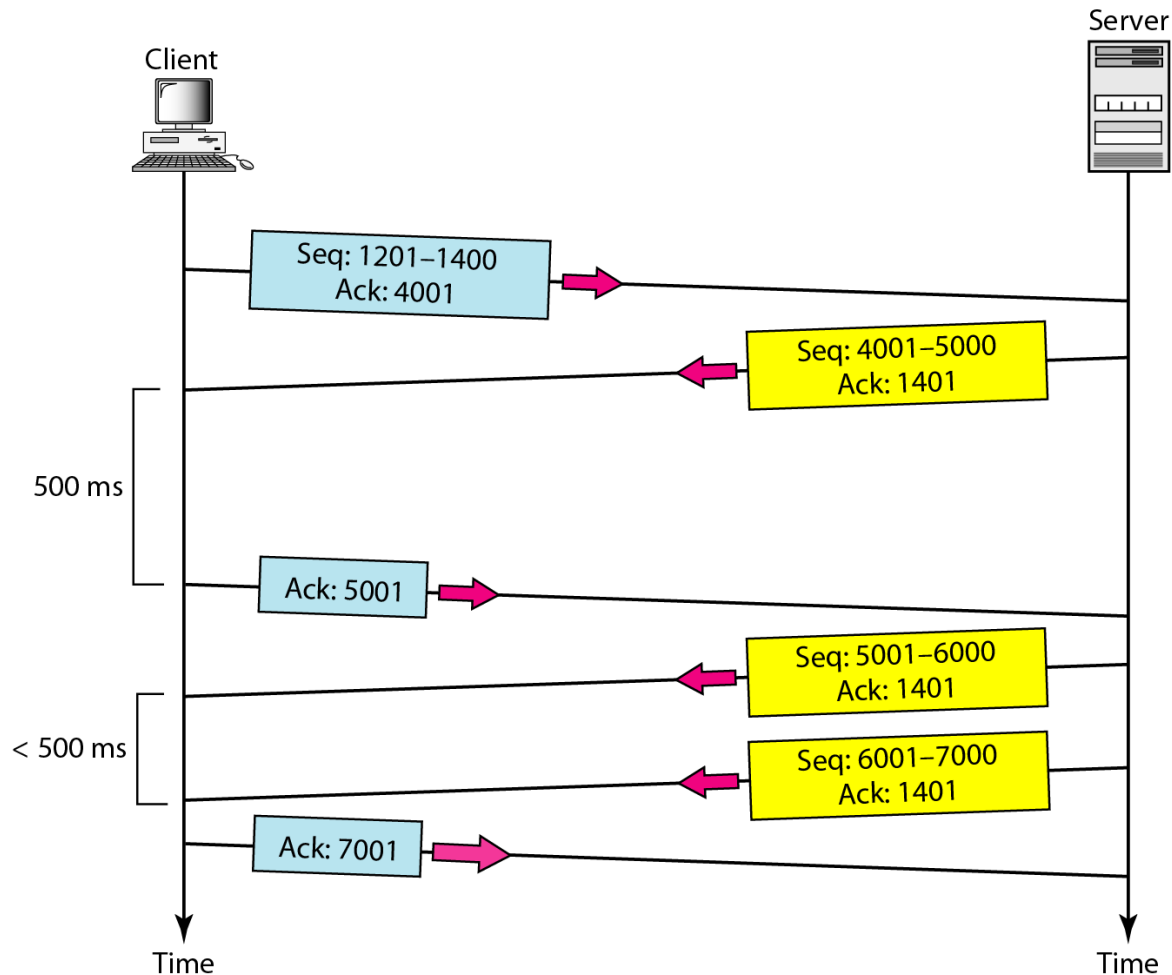
- ❑ Time-out period often relatively long:
 - long delay before resending lost packet
- ❑ Detect lost segments via duplicate ACKs.
 - Sender often sends many segments back-to-back
 - If segment is lost, there will likely be many duplicate ACKs.
- ❑ If sender receives for
the same data, it supposes that segment
was lost:
 - resend segment before timer
expires

TCP ACK generation [RFC 1122, RFC 2581]

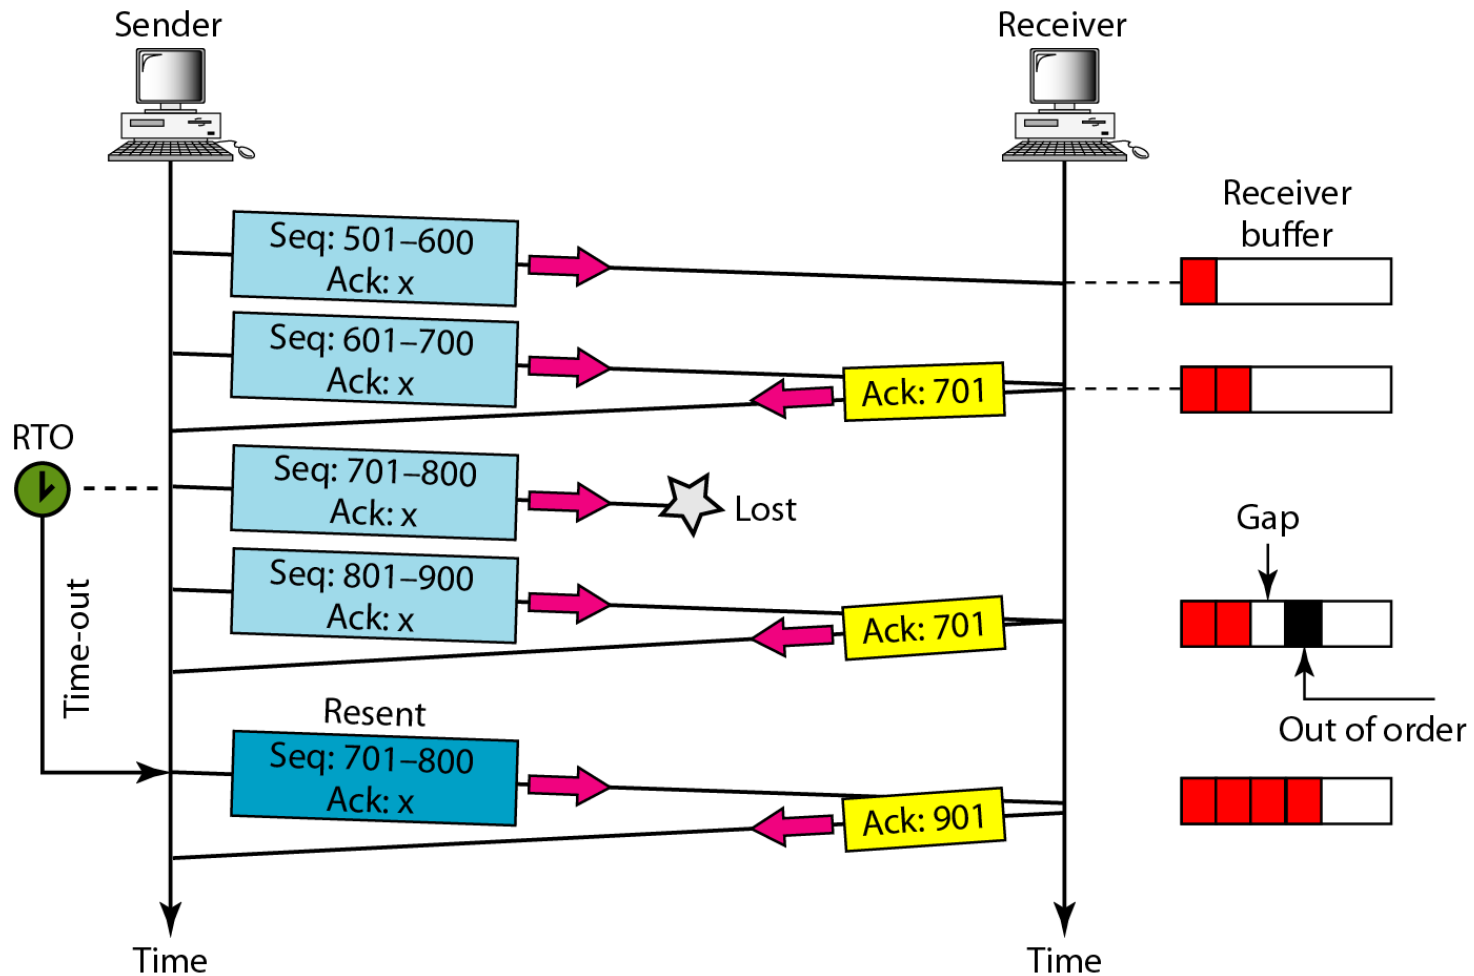
Event at Receiver	TCP Receiver action
Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	Delayed ACK. (within 500ms, typical value=100ms) If no next segment, send ACK
Arrival of in-order segment with expected seq #. One other segment has ACK pending	Immediately send single cumulative ACK, ACKing both in-order segments
Arrival of out-of-order segment higher-than-expect seq. # . Gap detected	Immediately send duplicate ACK, indicating seq. # of next expected byte
Arrival of segment that partially or completely fills gap	Immediate send ACK, provided that segment starts at lower end of gap

Delayed ACK: piggyback an ACK on data

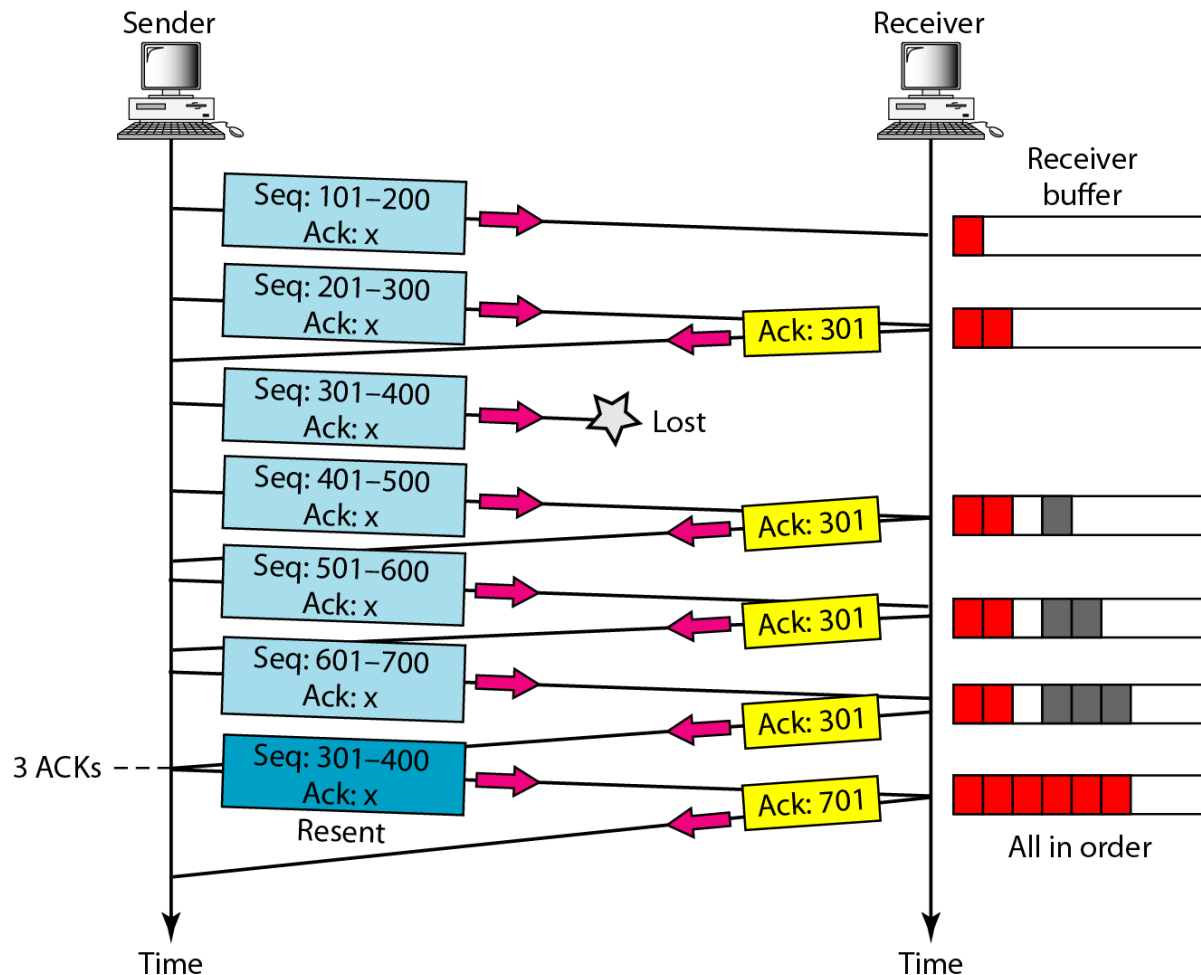
TCP Scenario: Normal case



TCP Scenario: Lost Segment

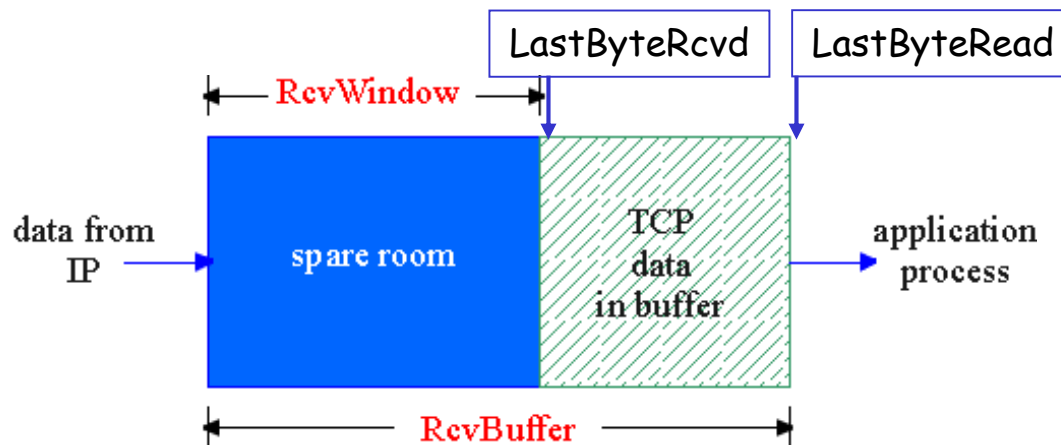


TCP Scenario: Fast Retransmission



TCP Flow Control

- ❑ To prevent the sender from overflowing the receiver's buffer.
- ❑ Dynamic window management
 - Receiver "advertises" its available buffer size in the "window size" field of ACK segments.
 - $\text{RcvWindow (or AdvertisedWindow)} = \text{RcvBuffer} - (\text{LastByteRcvd} - \text{LastByteRead})$
 - Sender limits unACKed data to RcvWindow
 - guarantees receive buffer doesn't overflow



TCP Flow Control

- ❑ If receiver buffer is totally full, then $RcvWindow = 0$
- ❑ Sender must entirely close its sender window and transmission may be halted forever!
 - Because TCP does not acknowledge ACK.
 - If the packet that opens the window is lost, the deadlock occurs!
- ❑ How to escape from this deadlock ?

TCP: Transmission Policy

❑ Sender Buffering

- 'Tinygram' wastes bandwidth
 - a keystroke in telnet session = 41 byte
(40 byte header + 1 byte data)
- be able to reduce header overhead by grouping many small data segments into one large TCP segment.
- **algorithm** (RFC 896)
when data come into the sender one byte, send the first byte. Then
 - 1) buffer all the rest until the outstanding byte is ACKed.
 - 2) Send a new packet if data fill the half the window or a maximum segment
 - better to be disabled if used on mouse movements.

TCP: Timers

- ❑ To perform its operation smoothly, TCP uses the 4 timers
 - Retransmission timer
 - Persistent timer
 - Keep-alive timer
 - 2MSL (time-waited) timer

- ❑ Retransmission timer
 - Usually TCP sender maintains one retransmission timer for each connection

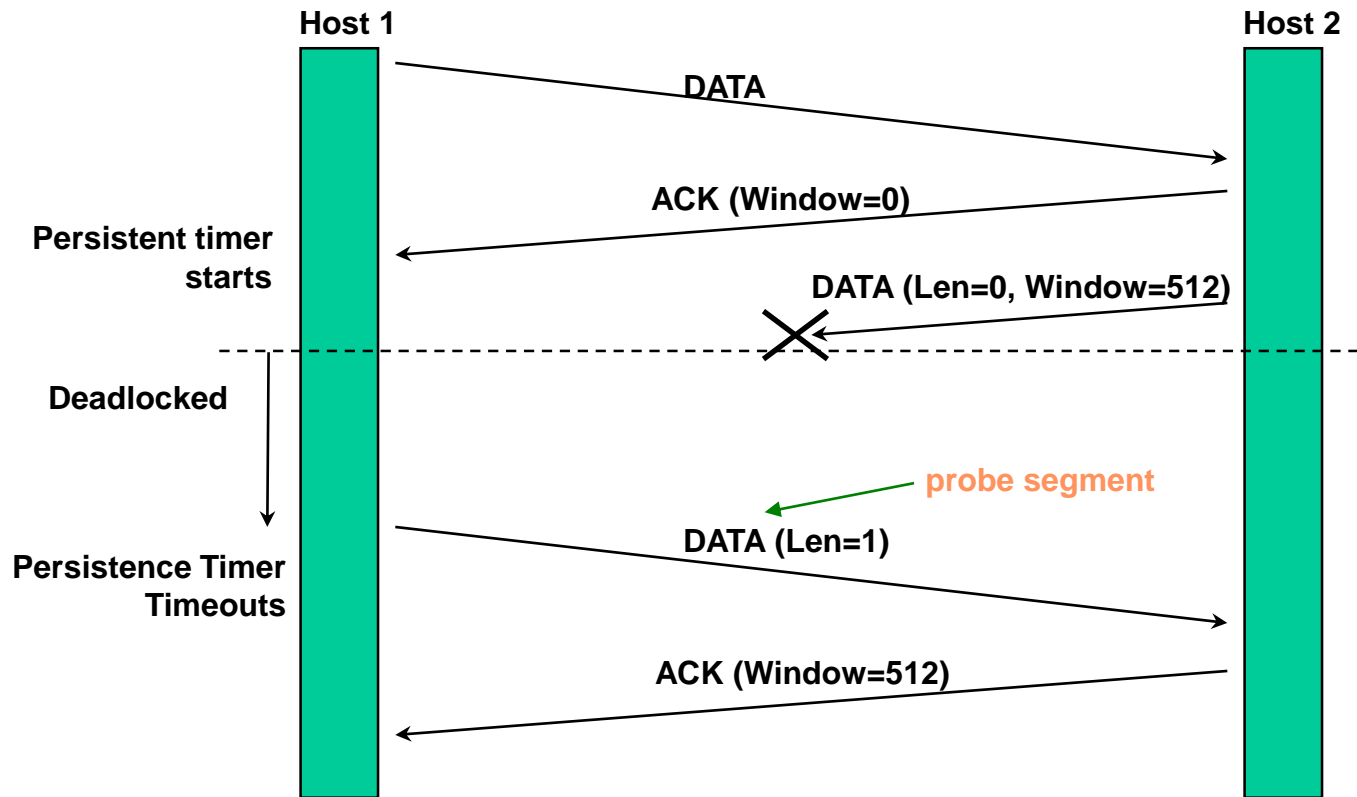
TCP: Timers

□ Persistent timer

- prevents deadlock from occurring when the packet with the update from the window size of 0 is lost.
- When the sending TCP receives an acknowledgment with a window size of zero, the persistence timer is started
- When persistence timer goes off, the sending TCP sends a special segment called a *probe segment*
- The probe alerts the receiving TCP that the acknowledgment was lost and should be resent.
- initially persistent timer = $2 \times \text{RTT}$;
- If a response is not received, the sender continues sending the probe segments and doubling and resetting the value of the persistence timer until the value reaches a threshold (usually 60 seconds).

TCP: Timers

□ Persistent timer



TCP: Timers

❑ Keep-alive timer

- Used to prevent a long idle connection between two TCPs.
- Each time the server hears from a client, it resets this timer.
- When timeout, send a packet to peer
 - usually timeout value = 2 hours
 - when timeout occurs, probe packets are sent every 75 sec.
 - If there is no response to 10 probe packets, the connection is torn down.

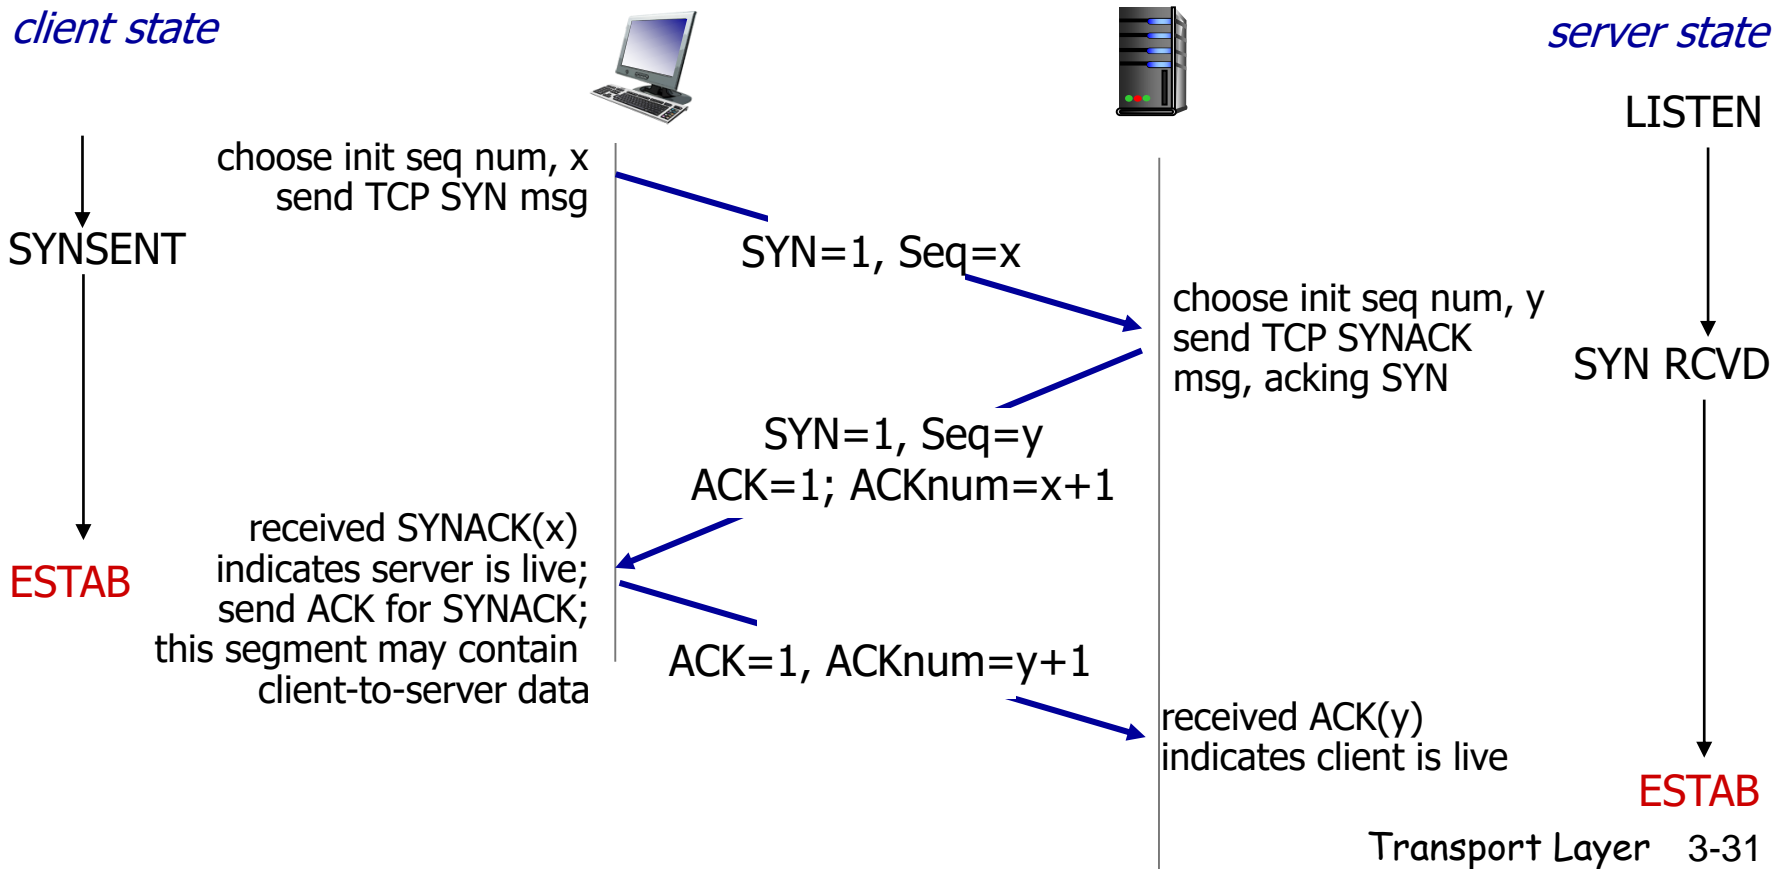
❑ 2MSL(maximum segment lifetime) timer

- Timer for the *TIMED WAIT* state while closing
 - (typically 120 sec.)
- To make sure that all the packets created by this connection have died off.

TCP Connection Management

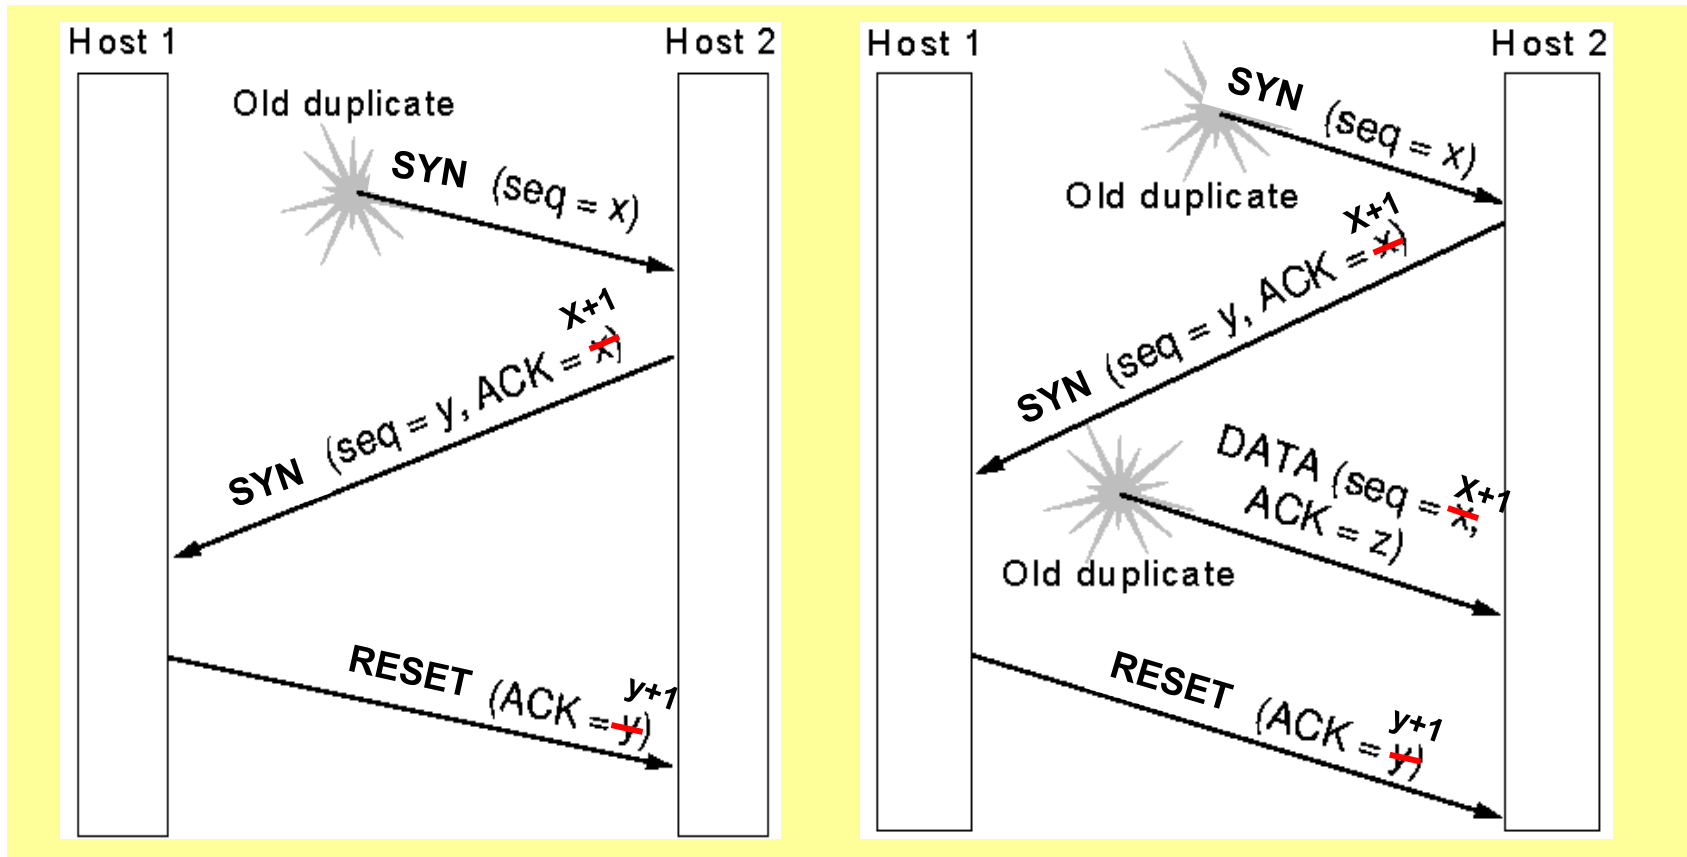
❑ Three-way handshake protocol

- A connection-oriented protocol establishes a virtual path between the source and destination before sending data.



TCP Connection Management

- ❑ Three-way handshake : against abnormal cases



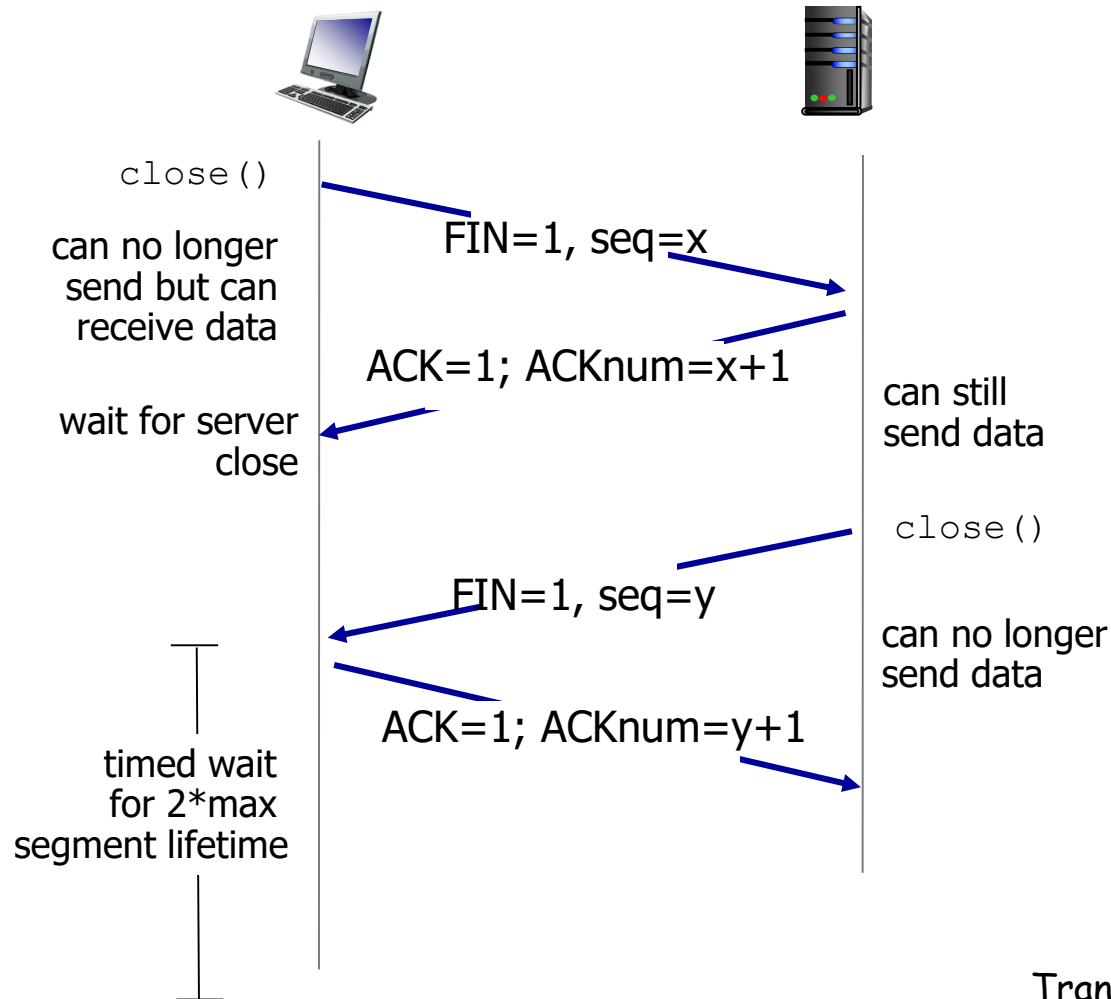
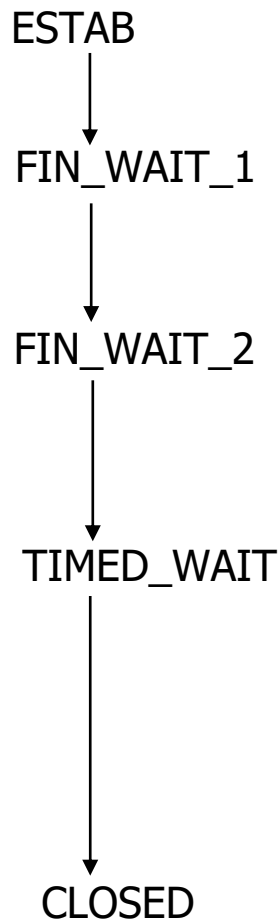
TCP: Closing a connection

- ❖ client, server each close their side of connection
 - send TCP segment with FIN bit = 1
- ❖ respond to received FIN with ACK
 - on receiving FIN, ACK can be combined with own FIN
- ❖ simultaneous FIN exchanges can be handled


TCP: Closing a connection

□ performed separately in each direction.

client state



server state

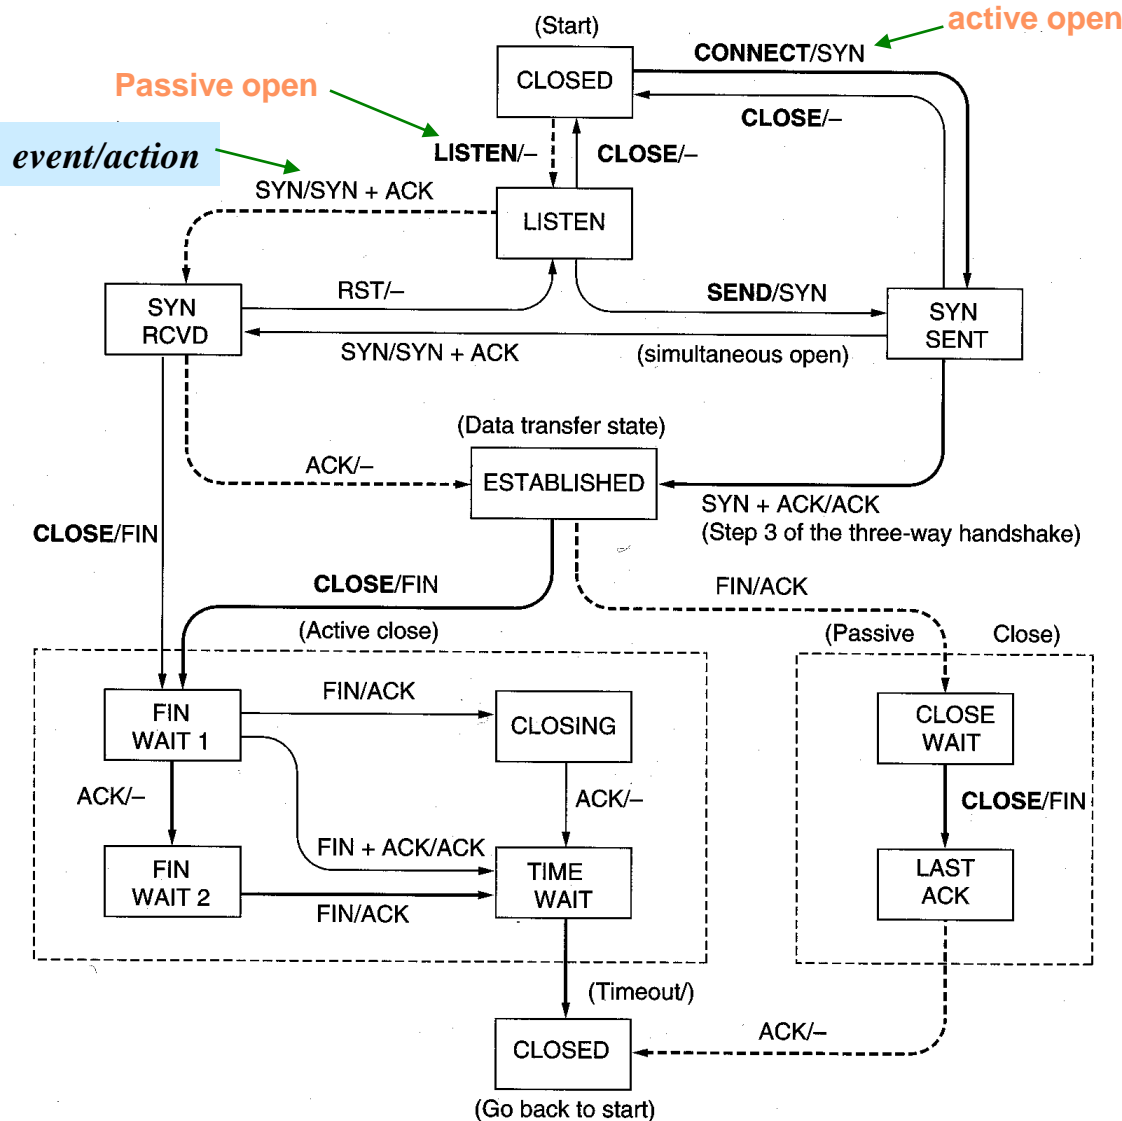


TCP: State Transition Diagram

dark line: client
dashed line: server

FIN WAIT 1: The client has said it is finished.
FIN WAIT 2: The server has agreed to release (half close)

TIME WAIT : Wait for all packets to die off.



TCP Connection Management

- ❑ 2MSL (Maximum Segment Lifetime) wait:
 - wait for final segment to be transmitted before releasing connection (typically 120 sec)
 - Socket *pair* cannot be reused during 2MSL
 - Delayed segments dropped
 - 2MSL effect
 - If you kill client and restart, it will get a different port
 - 2MSL wait protects against delayed segments from the previous “incarnation” of the connection.